



ICTSS00120 - Artificial Intelligence Skill Set

Session 2: Data for Machine Learning

Lecturer: Jordan Hill

Today's Agenda

- Introduction and Objectives
- IPython Environment and Kaggle
- Types of Data and Data Sources
- Understanding ML Attributes
- Relevant Algorithms for Datasets
- Data Preprocessing Techniques (Briefly)
- In-Class Activity: Introduction to Kaggle environment and Data Cleaning
- Summary and Q&A


Learning Objectives

- Understand the IPython environment and how to use Kaggle.
- Identify key data sources for machine learning.
- Differentiate between structured, unstructured, labelled, and unlabelled data.
- Comprehend ML attributes: data attributes, model attributes, and target attributes.
- Recognize relevant algorithms for different types of datasets.
- Grasp data biases and methods for generating balanced data.
- Introduction to data preprocessing techniques such as randomizing, deduplication, and data cleaning.
- Understand the importance of creating training, validation, and test subsets.

Introduction to IPython and Kaggle

IP[y]: IPython Interactive Computing

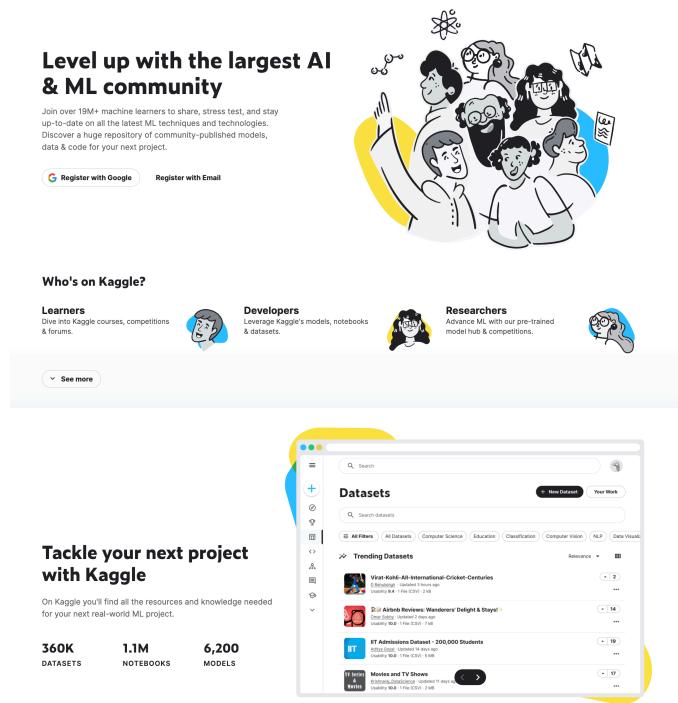
What is it? And why do we use it??



- Interactive computing environment for writing and sharing Python code.
- Supports visualization, markdown, and LaTeX.
- **Why Important?**
 - Facilitates exploratory data analysis.
 - Enhances code readability and sharing.
 - Enables reproducibility of experimentation results



kaggle



- Online platform for data science competitions and collaboration.
- Provides datasets, code snippets, and a notebook environment.
- Provides GPU accelerated environment for working with vision and LLM tasks!

Key Data Sources

Identifying Key Data Sources:

- **Public Repositories:** Examples include UCI ML Repository, Kaggle Datasets, and government open data websites.
- **APIs:** Data from web services such as Twitter, OpenWeatherMap, etc.
- **Internal Data:** Company databases, CRM systems, and internal APIs.



Types of Data

Structured Data:

- Clearly defined data types.
- Organized in rows and columns.
- Examples: Databases, spreadsheets.

Unstructured Data:


- No predefined structure.
- Text, images, videos, etc...
- Examples: Emails, social media posts.

UNSTRUCTURED DATA



VS

STRUCTURED DATA



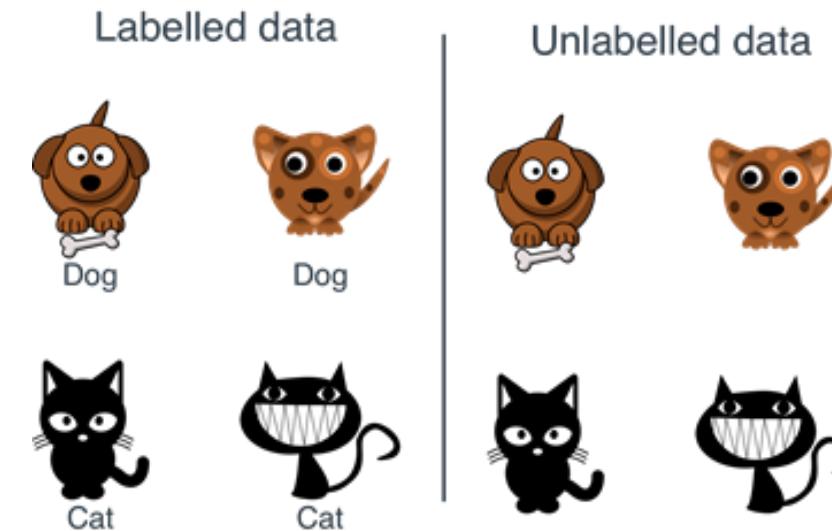
Types of Data (Machine Learning / AI)

Labelled Data:

- Each data point is assigned a label.
- Used for supervised learning.
- Example: Emails labelled as ‘spam’ or ‘not spam’.

Unlabelled Data:

- No labels are assigned to data points.
- Used for unsupervised learning.
- Example: Customer purchase data without any categorization.



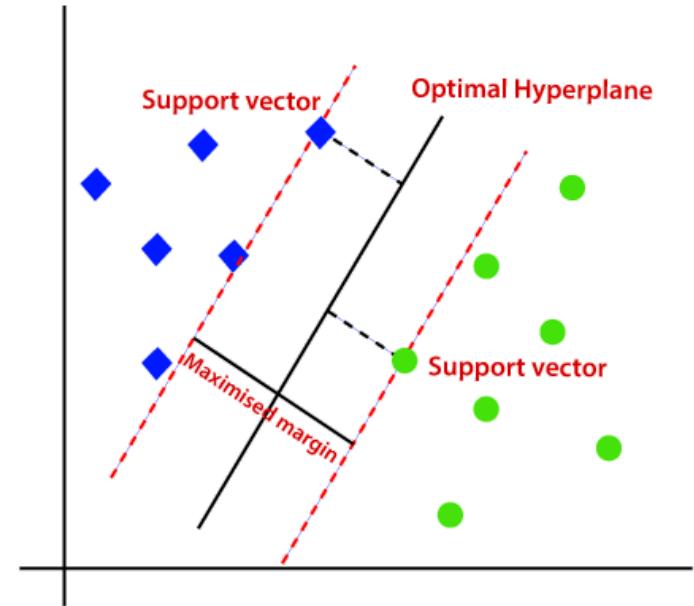
Relevant Algorithms for Labelled & Unlabelled Data

Labelled Datasets:

- **Supervised Learning:**
 - Regression: Linear Regression. [Data Regression](#)
 - Classification: Decision Trees, SVM, Naive Bayes, KNN. [Classification in ML](#)

Unlabelled Datasets:

- **Unsupervised Learning:**
 - Clustering: K-Means, DBSCAN. [What is K-Means?](#)
 - Dimensionality Reduction: PCA, LDA. [What is PCA?](#)



Understanding ML Attributes

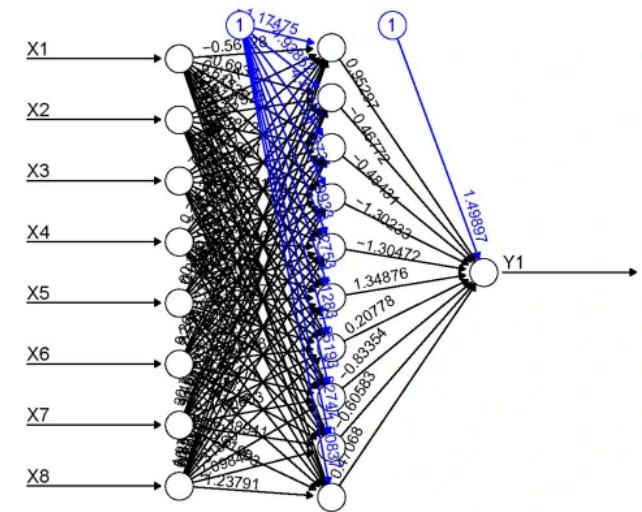
ML Attributes:

- **➡ Data Attributes:** Features or predictors used to train a model.
- **Model Attributes:** Parameters of the model.
- **Target Attributes:** The outcome variable the model is predicting.

Understanding ML Attributes

ML Attributes:

- **Data Attributes:** Features or predictors used to train a model.
- **→ Model Attributes:** Parameters of the model.
- **Target Attributes:** The outcome variable the model is predicting.



Understanding ML Attributes

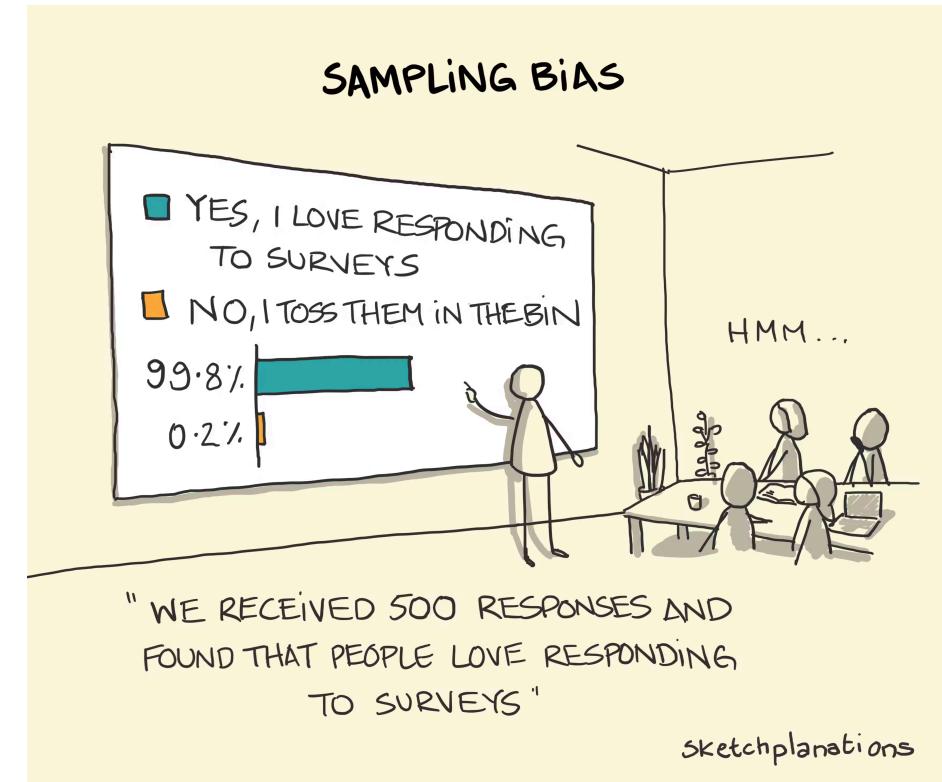
ML Attributes:

- **Data Attributes:** Features or predictors used to train a model.
- **Model Attributes:** Parameters of the model.
- ➔ **Target Attributes:** The outcome variable the model is predicting.

Understanding Data Biases

Statistical Bias:

- **Definition:** Systematic error introduced into data collection.
- **Examples:**
 - Sampling bias.
 - Measurement bias.
- **Mitigation Techniques:**
 - Ensure diverse dataset collection.
 - Resampling or reweighting data.



Bias is a broader conversation when it comes to AI systems

1. Cognitive Bias — Bias fed by limitations in our own cognition & perspective
2. Statistical Bias — Affects the quality of our models
3. Prejudice / Descrimination — Prejudice on the basis of race, sex, age, etc...
4. Conflicts of Interest — Corruption / vested interests

The world may be biased in ways we don't want to perpetuate!

We want to design AI systems that are benevolent and beneficial.

As data scientists, we have a moral obligation to ensure we develop systems ethically. Notions of **good & bad, right & wrong** are highly relevant to the design of AI systems.

Question your Data!

How can we actually identify & deal with *statistical* bias?



Identifying and Reducing Statistical Bias

Methods for Identifying Statistical Bias

- **Exploratory Data Analysis (EDA):**
 - Use visualizations (e.g., histograms, box plots) to detect distributions and outliers.
 - Summarize data using descriptive statistics (mean, median, mode).
- **Statistical Tests:**
 - Perform tests (e.g., Chi-square test, T-test) to check for bias in dataset distributions.
- **Correlation Analysis:**
 - Evaluate correlations between features to identify multicollinearity issues.
- **Audit for Missing Data:**
 - Identify patterns in missing data to detect if data is missing at random (MAR) or not at random (MNAR).

Generating Balanced Data

Generating Balanced Data:

- **Techniques to Balance Data:**
 - Oversampling minority class.
 - Undersampling majority class.
 - SMOTE (Synthetic Minority Over-sampling Technique).

Reducing Statistical Bias in Datasets

1. Data Collection/Preperation Improvements:

- Use random sampling of larger populations to avoid an **introduced** sampling bias when collecting data.
- Ensure diverse and representative dataset collection.

2. Data Resampling Techniques:

- **Oversampling:** Increase the number of minority class samples.
 - Example: [SMOTE \(Synthetic Minority Over-sampling Technique\)](#).
- **Undersampling:** Reduce the number of majority class samples.
 - Example: [Random undersampling](#).

Reducing Statistical Bias in Datasets

3. Data Augmentation:

- Create synthetic data points to balance datasets.
- Techniques include geometric transformations for images, paraphrasing for text.

4. Normalization and Standardization:

- Scale numerical features to have similar ranges, reducing the impact of features with larger scales.
- Normalize data to have a mean of zero and a standard deviation of one.

Reducing Statistical Bias in Datasets

5. Cross-Validation:

- Use K-fold cross-validation to ensure bias isn't introduced through data splitting.
- Helps in assessing model performance on different subsets of data.

6. Bias Correction Algorithms:

- Implement algorithms designed to reduce bias, such as reweighting or adversarial de-biasing.

The Importance of Balance in Data

- Imbalanced datasets can lead to biased models that favor majority classes.
- Balancing the data ensures that the model treats all classes equally and performs better on minority classes.

Techniques to Balance Data:

1. Oversampling

Definition: Increasing the number of samples in the minority class to match the majority class.

2. Undersampling

Definition: Reducing the number of samples in the majority class to match the minority class.

3. Data Augmentation

Definition: Creating new data points to add to the dataset, often used for image or text data.



Key Techniques Explained

SMOTE (Synthetic Minority Over-sampling Technique)

- **How It Works:**
 - For each minority class sample, find its k-nearest neighbors.
 - Randomly select one of the neighbors and create a synthetic sample along the line between the original sample and its neighbor.

Random Undersampling

- **How It Works:**
 - Randomly remove samples from the majority class.
 - This reduces the number of majority class samples to match the number of minority class samples.

Data Augmentation for Images

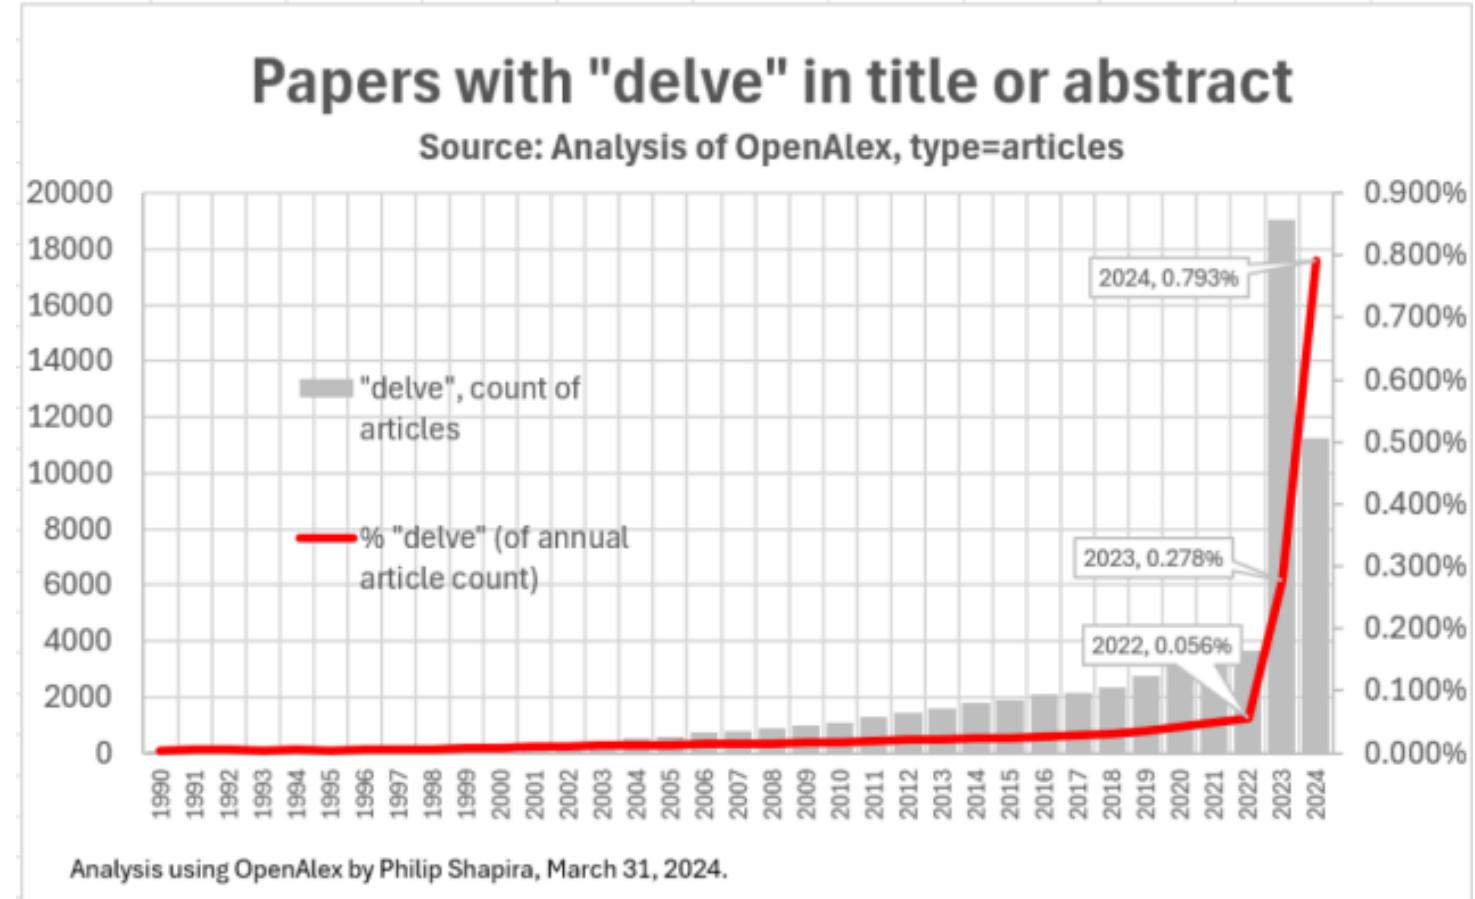
- **How It Works:**
 - Apply transformations such as rotations, flips, crops, and color adjustments to existing images.
 - This generates new, varied images while preserving labels.

Data Augmentation for Text

- **How It Works:**
 - Generate new sentences with the same meaning but different structure.
 - Use techniques such as paraphrasing, synonym replacement, or back-translation.

It seems like OpenAI forgot to do this last one with "delve",
ChatGPT seems to really like it for some reason...

Whoops!



Data Preprocessing for ML

Randomizing: Shuffle data to prevent order effects.

Deduplication: Remove duplicate records.

Data Cleaning: Handle missing values, correct inconsistencies.



DATA PREPROCESSING

Steps for Data Preprocessing



Creating Data Subsets

Creating Data Subsets:

- **Training Subset:** Used to train the model.
- **Validation Subset:** Used to tune model parameters.
- **Test Subset:** Used to evaluate the final model performance.

[Good Overview Wikipedia](#)

Why Important?

- Ensures the model generalizes well to unseen data.
- Helps in detecting overfitting and underfitting.

In-Class Activity

Activity:

- Load a dataset into Kaggle.
- Evaluate its suitability for machine learning.
- Visualize the dataset and perform initial exploratory data analysis.

Lab Sheet

1. Log in to Kaggle.
2. Find a dataset.
3. Load it into a notebook.
4. Visualize basic statistics and charts. (ChatGPT/Co-pilot can help with this!)

Questions to Explore & Discuss with your Peers:

- What kind of data (structured/unstructured; labelled/unlabelled) is this?
- What preprocessing steps might be necessary?
- What biases might exist in the data?
- Are there any biases in the data? How could you identify them?
- What could be done to correct bias?
- How would you split the data into training, validation, and test sets?

Summary

- Introduced to IPython, Kaggle, and key data sources.
- Discussed types of data and ML attributes.
- Covered relevant algorithms for different datasets.
- Understood data biases and techniques for balanced data.
- Learned data preprocessing steps and importance of data subsets.

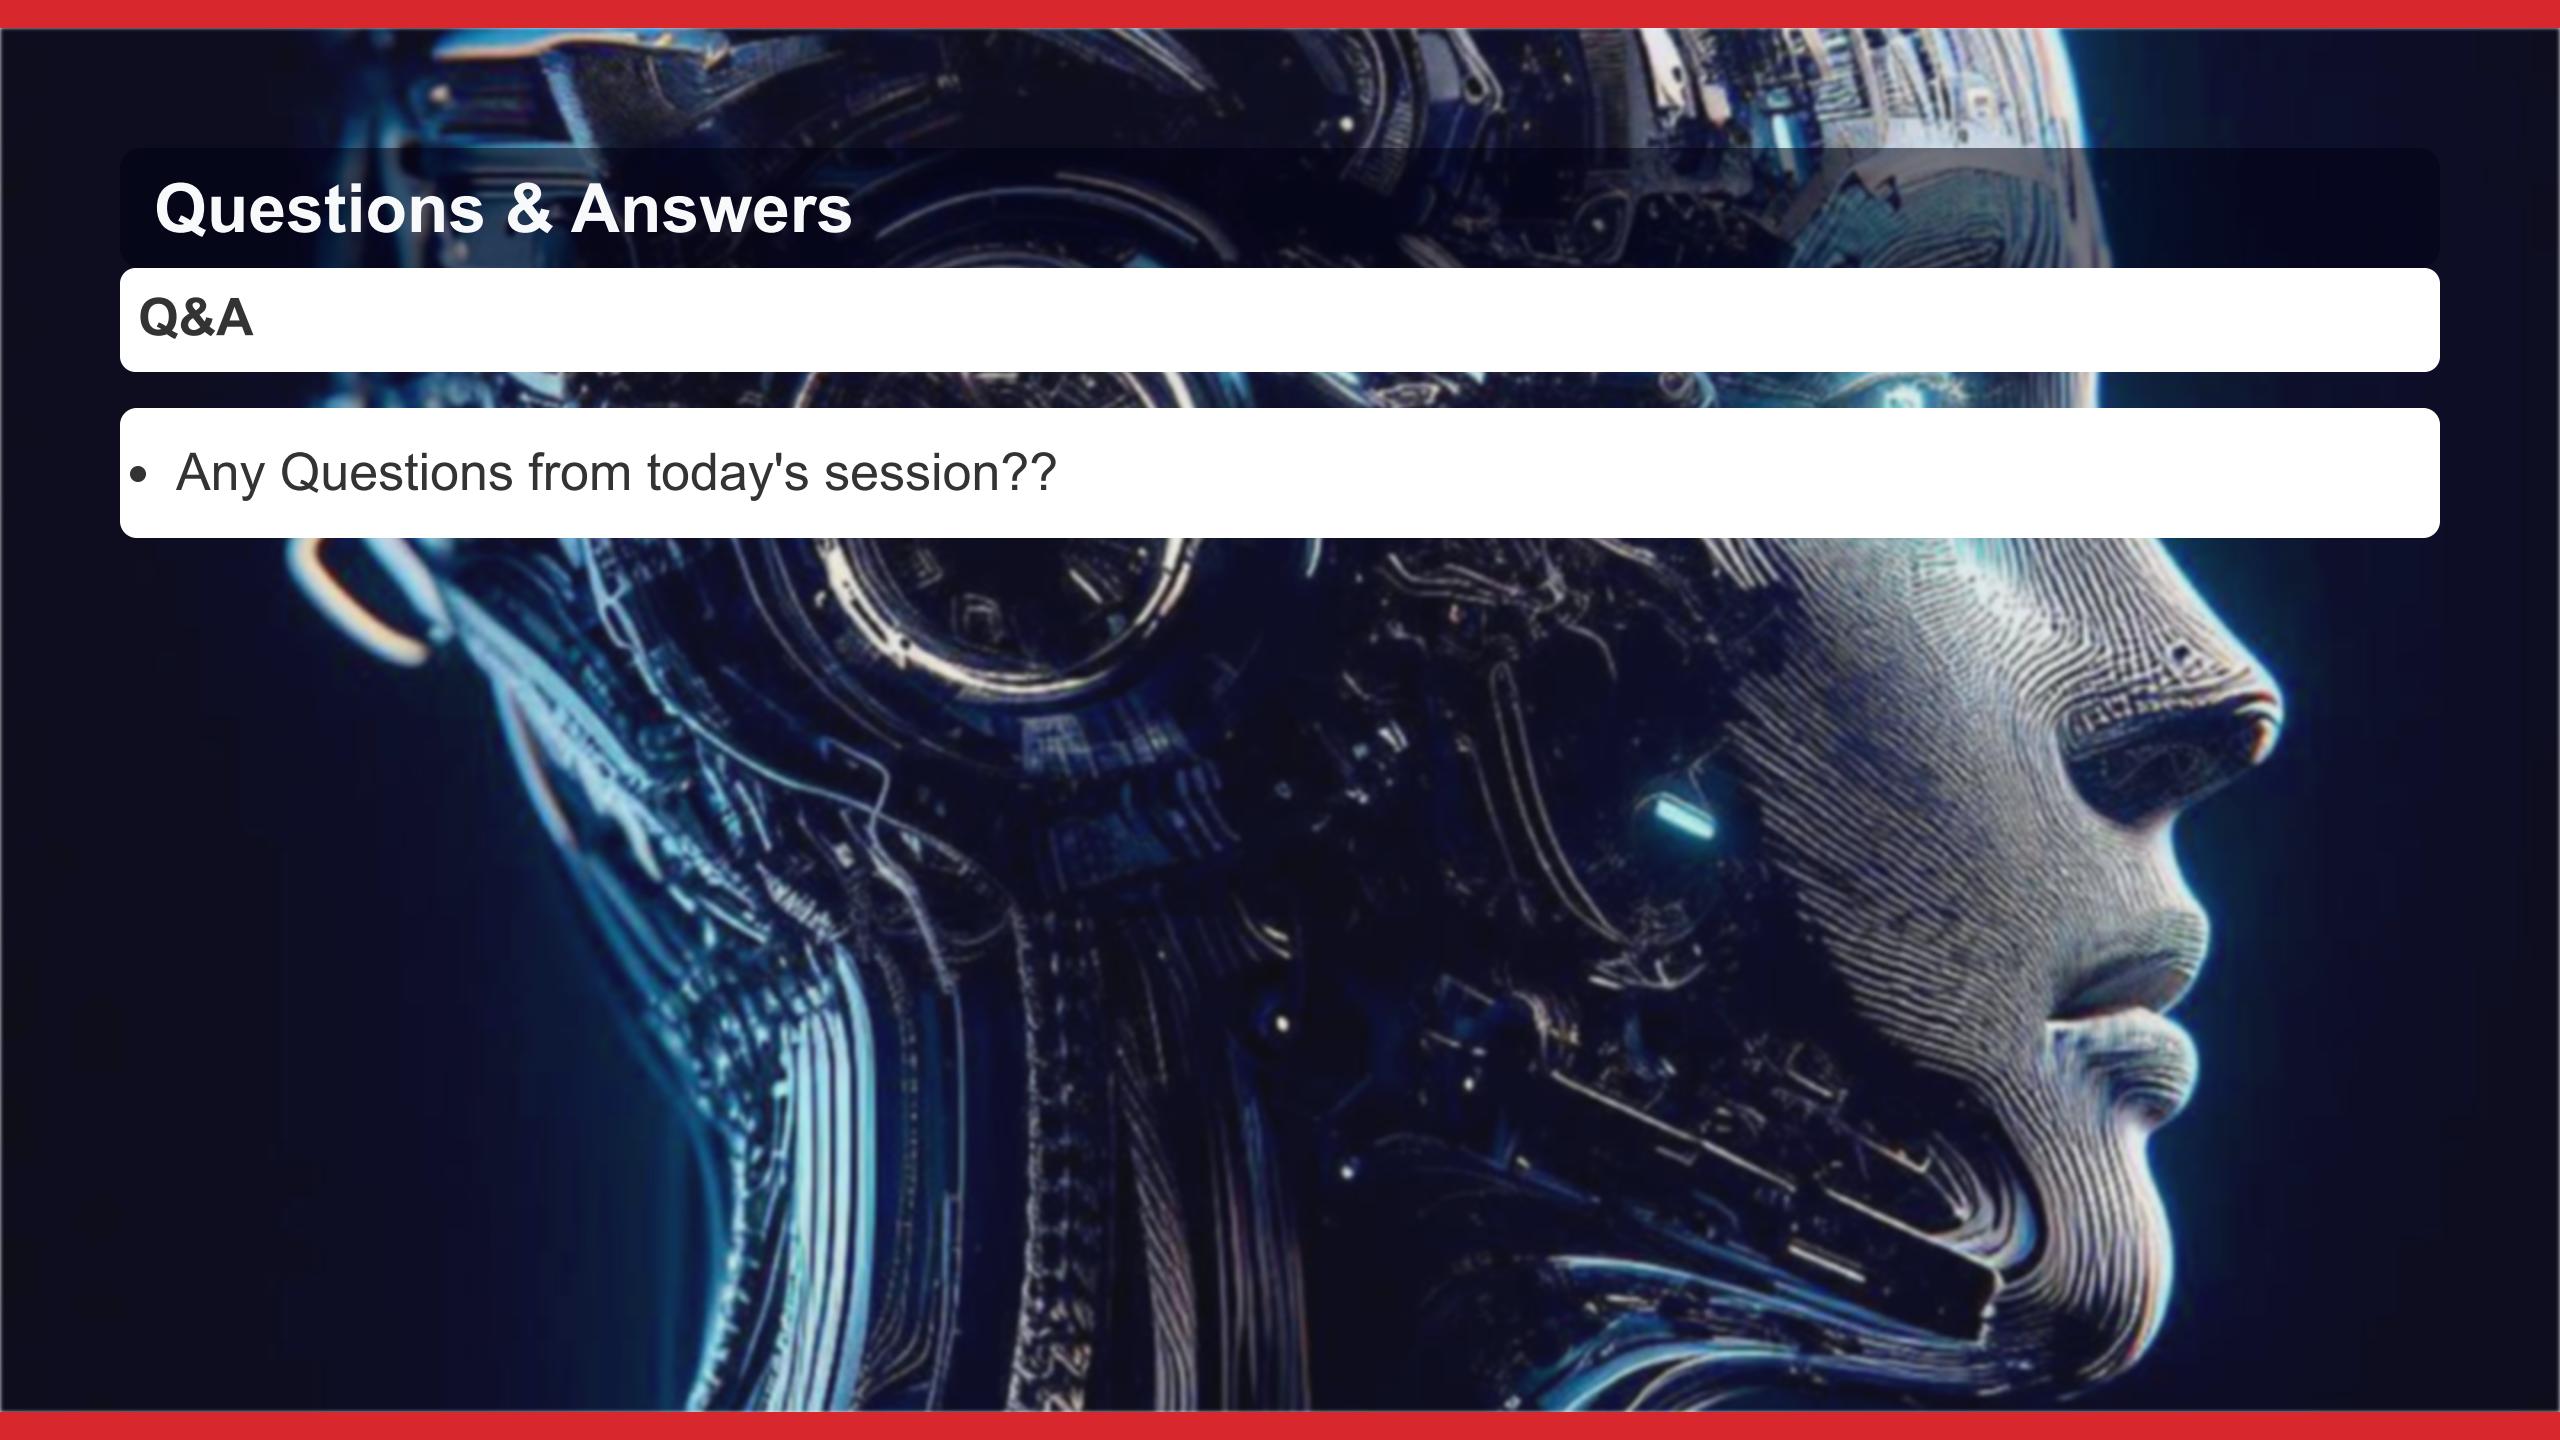
Key Takeaways

- **IPython and Kaggle:** Essential tools for data manipulation and collaboration.
- **Data Types:** Understanding structured vs. unstructured and labelled vs. unlabelled data.
- **ML Attributes:** Importance of data attributes, model attributes, and target attributes.
- **Algorithms:** Supervised and unsupervised learning techniques.
- **Bias and Balancing:** Techniques to identify and mitigate biases in datasets.
- **Preprocessing:** Critical steps in preparing data for machine learning.

Questions & Answers

Q&A

- Any Questions from today's session??



Next Week

Next Week:

- Essentials of ML and DL Technologies.
- Dive deeper into supervised and unsupervised learning techniques.
- Intro to sci-kit learn