

ICTSS00120 - Artificial Intelligence Skill Set

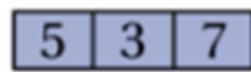
Session 10: Fundamentals in Deep Learning & Introduction to Machine Vision

Lecturer: Jordan Hill

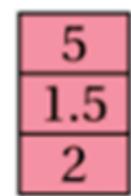
Learning Objectives

- Understand the basics of linear algebra and tensors for deep learning.
- Explore computing tensors, GPU vs CPU.
- Learn feature engineering and data preprocessing techniques for deep learning.
- Understand activation functions and their roles.
- Delve deeper into Convolutional Neural Networks (CNNs) and their applications in machine vision.

(11)



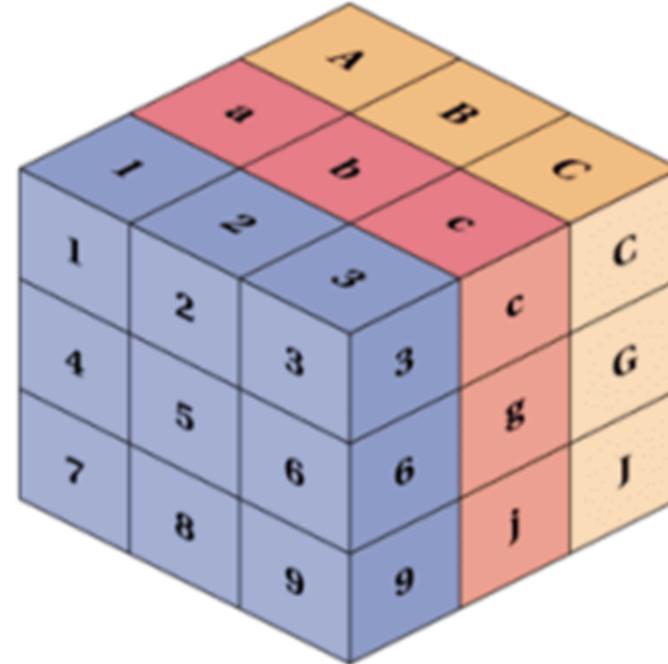
SCALAR



Column Vector
(shape 3x1)



MATRIX



TENSOR

Understanding Scalar, Vectors, Matrices and Tensors

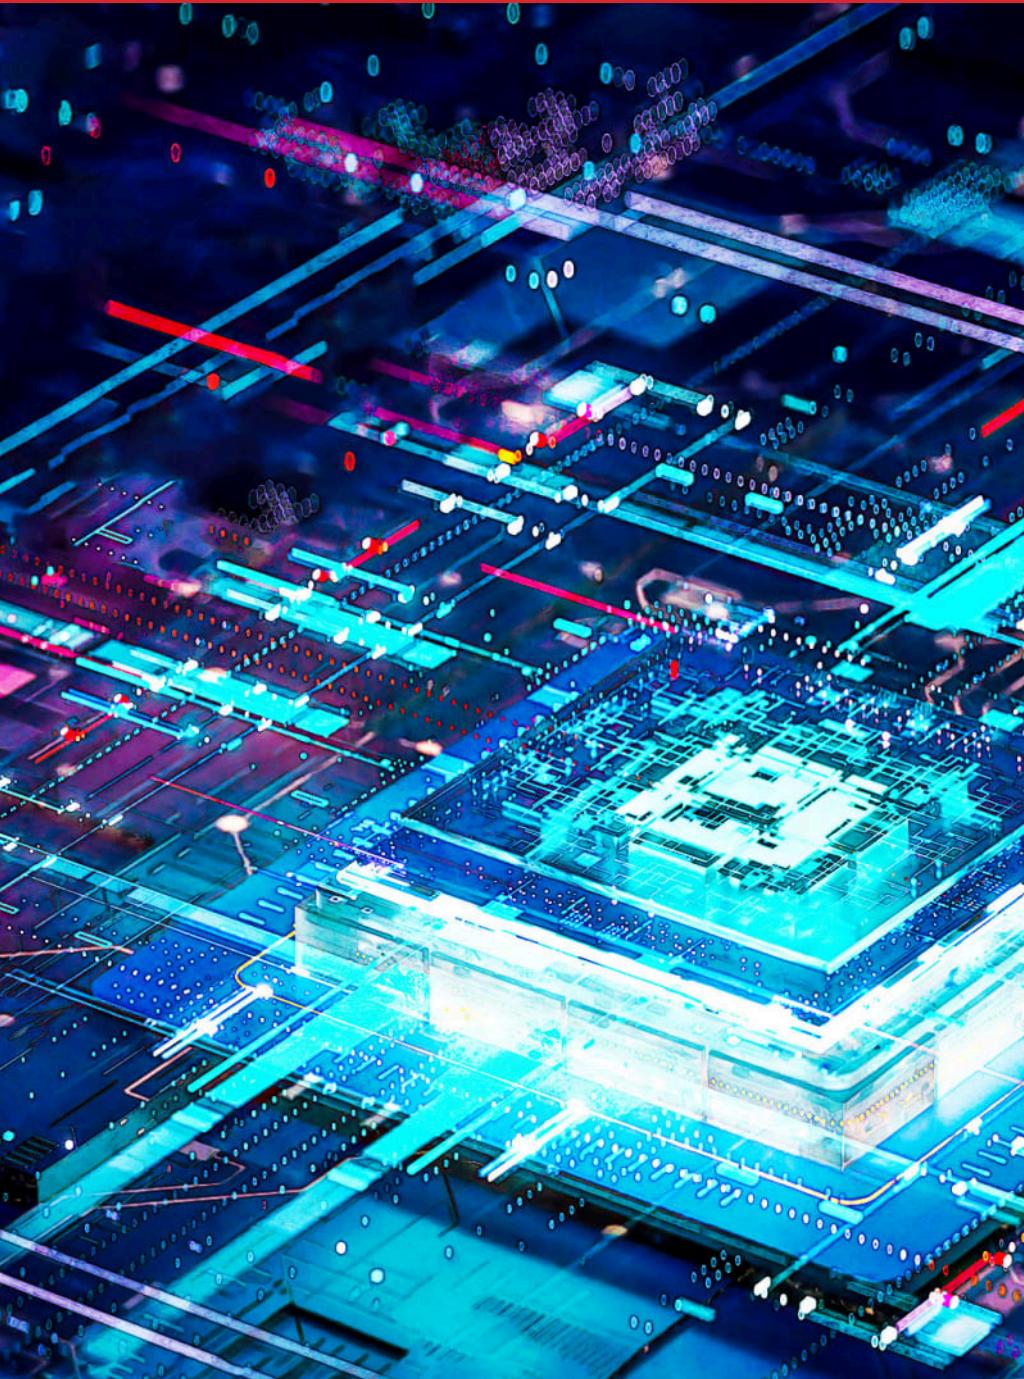
Linear Algebra for Machine Learning

Scalars, Vectors, and Tensors

- **Scalar:** A single number.
 - Example: $a = 1$
- **Vector:** An array of numbers.
 - Example: $\mathbf{v} = [1, 2, 3]$
- **Matrix:** A 2D array of numbers.
 - Example: $\mathbf{M} = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$
- **Tensor:** An n-dimensional array of numbers.
 - Example: \mathbf{T} could be a 3D array, a 4D array, etc.

Computing Tensors: GPU vs CPU

- **Central Processing Unit (CPU)**: General-purpose processor.
 - Best for tasks with lower parallelism.
- **Graphics Processing Unit (GPU)**: Specialized for highly parallel tasks.
 - Crucial for deep learning, handling large tensors efficiently.



Computing Tensors: TPU vs NPU

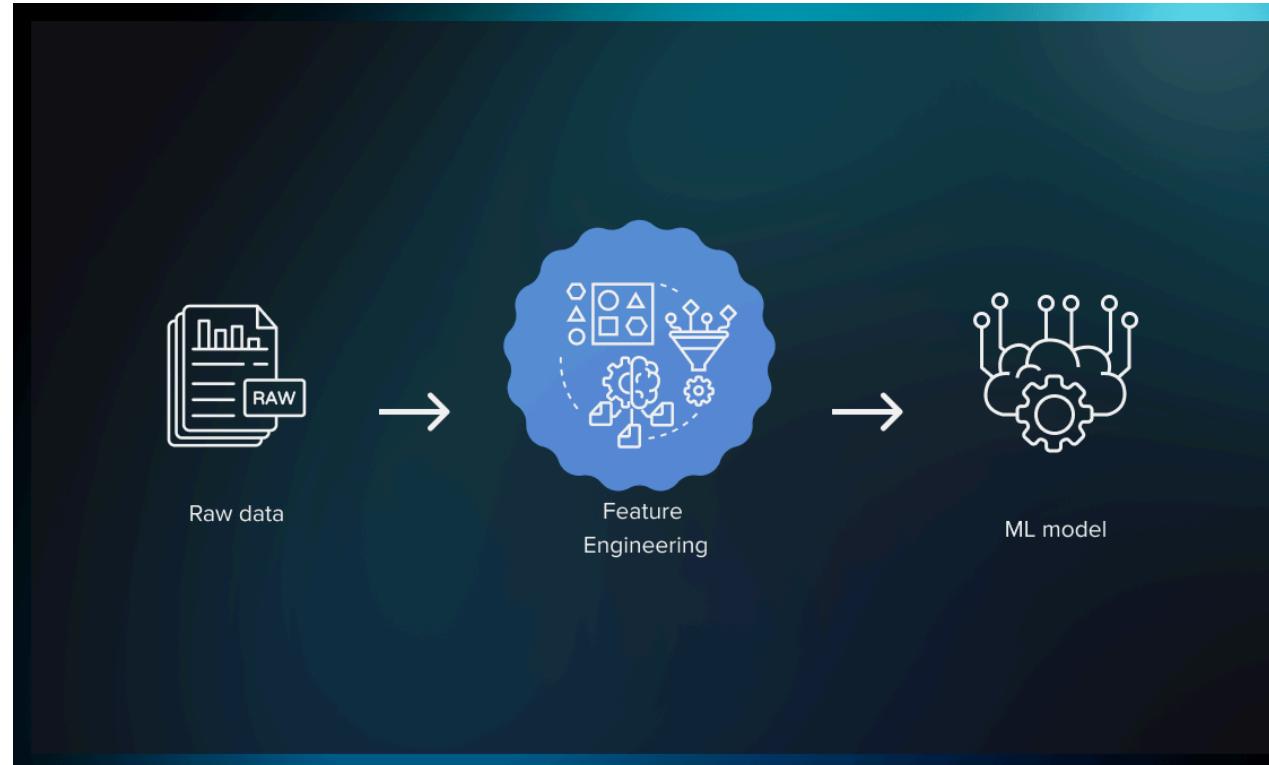
- **NPU (Neural Processing Unit):**
 - **Architecture:** Specialized for neural network computations.
 - **Applications:** Real-time applications like translation, facial recognition, voice assistants.
- **TPU (Tensor Processing Unit):** (Google)
 - **Architecture:** Optimized for TensorFlow operations.
 - **Applications:** Data centers, training large machine learning models, Google Cloud services.



Feature Engineering and Data Preprocessing

Techniques for Deep Learning

- **Normalization:** Scaling features to a standard range.
- **Standardization:** Transforming data to have zero mean and unit variance.
- **Data Augmentation:** Creating new training instances by modifying existing data (e.g., rotations, flips for images).



Activation Functions

Key Functions

- **ReLU (Rectified Linear Unit):**

$$f(x) = \max(0, x)$$



- Benefits: Avoids vanishing gradients, computationally efficient.

- **Sigmoid:** $f(x) = \frac{1}{1+e^{-x}}$

- Benefits: Outputs in range (0, 1), used for probability estimation.

- **Tanh:** $f(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$

- Benefits: Outputs in range (-1, 1), zero-centered.



Convolutional Neural Networks (CNNs)

Basics and Applications

- **Structure:**
 - **Convolutional Layers:** Apply convolution operations to extract features.
 - **Pooling Layers:** Reduce the dimensionality of feature maps.
 - **Fully Connected Layers:** Perform high-level reasoning and classification.
- **Applications:**
 - Image recognition, object detection, and video analytics.



Deep Dive into CNNs

Understanding Convolutional Layers

- **Convolutions**: Apply filters (kernels) to the input image to create feature maps.
- **Filters**: Small matrices that slide over the input to detect patterns (e.g., edges, textures).



Deep Dive into CNNs

Pooling Layers


- **Max Pooling**: Takes the maximum value in a window (e.g., 2x2) to down-sample the feature map.
- **Average Pooling**: Takes the average value in a window to down-sample.

Practical Applications of CNNs

Image Recognition and Beyond

- **Image Classification**: Recognizing objects in images (e.g., cats, dogs, cars).
 - Example: AlexNet winning the ImageNet competition.
- **Object Detection**: Identifying and locating objects within images.
 - Example: YOLO (You Only Look Once) detecting multiple objects in real-time.
- **Segmentation**: Dividing an image into meaningful regions.
 - Example: U-Net for medical image segmentation.





Lab: Building a Simple CNN for Image Classification

Let's look at CNN and Machine vision more closely

Lab Sheet

Summary and Q&A

Summary:

- Learned about basics of linear algebra and tensors.
- Explored GPU vs CPU for computing tensors.
- Learned feature engineering and data preprocessing techniques.
- Understood key activation functions.
- Delved deeper into CNNs and their applications.

Q&A:

- Any questions about today's topics?

Contact: jordan.hill@nmtafe.wa.edu.au

Homework

Next Week:

- Introduction to Recurrent Neural Networks (RNNs).
- Explore RNNs and their applications in sequence prediction.

Tasks:

1. Review key concepts covered today.
2. Read about RNNs and their applications.
3. Watch related videos on deep learning and CNNs.